Data Integration Log — HeavySuppliersWarehouseDatasets

Built on top of the cleaned files (see CLEANING_LOG.md) — three analysis-ready, denormalized views:

1. sales_transactions.csv (130,402 rows — line-item grain)

`sales_orders_lines` LEFT JOIN `sales_orders_header` (on so_id) LEFT JOIN `customers` (on customer_id) LEFT JOIN `branches` (on branch_id) LEFT JOIN `products` (on product_id) LEFT JOIN `invoices` (on so_id) - Row count matches sales_orders_lines exactly — verified no fan-out. - 12,746 rows have no invoice_id: these are line items on Cancelled orders. Confirmed invoices only ever exist for Delivered orders (checked before merging), so these nulls are expected, not a defect.

2. purchase_transactions.csv (155,495 rows — line-item grain)

`purchase_orders_lines` LEFT JOIN `purchase_orders_header` (on po_id) LEFT JOIN `suppliers` (on supplier_id) LEFT JOIN `branches` (on branch_id) LEFT JOIN `products` (on product_id) - Row count matches purchase_orders_lines exactly.

3. inventory_position.csv (180 rows — one per product/branch)

`inventory_master` LEFT JOIN `products` (on product_id) LEFT JOIN `branches` (on branch_id) - Added `stock_value` = current_stock x unit_cost. - Carries forward the `scale_mismatch_flag` from cleaning, so anyone using stock_value downstream is warned the current_stock scale is still unresolved against this table’s own reorder/max fields.

Why these three and not one giant table

Sales and purchases have different grains, dimensions, and business meaning (revenue vs. cost side) — merging them into one table would force either a many-to-many join or duplicate dimension columns with no shared fact. Inventory is a snapshot, not a transaction, so it stays separate. Each file is ready to pivot/group directly in Excel, Python, or BI tools without further joining.

Not merged in

stock_ledger (237K movement rows) and payments (19K rows) were left out of these three files by design — they don’t join 1:1 onto a single line item (a payment can be ambiguous per the cleaning log; a stock movement is its own timeline). They’re better analyzed as their own tables, joined in only for a specific question.